

1. Download the data from [here](#)
2. Code the model to classify data like below image



3. Write your own callback function, that has to print the micro F1 score and AUC score after each epoch.
4. Save your model at every epoch if your validation accuracy is improved from previous epoch.
5. you have to decay learning based on below conditions.
- Cond1. If your validation accuracy at that epoch is less than previous epoch accuracy, you have to decrease the learning rate by 0.5, if not, keep it as it is.
 - Cond2. For every 3rd epoch, decay your learning rate by 5%.
6. If you are getting any NaN values(either weights or loss) while training, you have to terminate your training.
7. You have to stop the training if your validation accuracy is not increased in last 2 epochs.
8. Use tensorboard for every model and analyse your gradients. (you need to upload the screenshots for each model for evaluation)
9. use cross entropy as loss function
10. Try the architecture params as given below.

```
In [11]: from google.colab import drive
drive.mount('/content/drive')

In [59]: import tensorflow as tf
import datetime
import numpy as np
import pandas as pd
from tensorflow.keras.callbacks import EarlyStopping,
filterwarnings(action='ignore', module='tensorflow')
filterwarnings(action='ignore', module='tensorboard')

In [60]: # Point 1
data = pd.read_csv('/content/drive/MyDrive/Calibacks/data.csv')
print (data.shape)
data.info()

Out[60]: (28000, 3)
         f2  label
0  0.450564  1.04305    0
1  0.005832  0.967862   10
2  0.011726  0.971021   10
3  0.002179  0.380488   00
4  0.702552  0.955880   00

In [61]: y = data['label'].values
data.drop(['label'], axis=1)
```

```
In [68]: x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.33, stratify=y)

In [70]: print (X_train.shape)
print (X_test.shape)

(13400, 2)
(6600, 2)

In [71]: validation_data=(X_test,y_test)
train_data=(X_train,y_train)
```

```
Model-1
1. Use tanh as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use RandomUniform(0,1) as initializer.
3. Analyze your output and training process.

In [72]: import tensorflow.keras.backend as K
from tensorflow.keras.callbacks import Callback
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.callbacks import LearningRateScheduler
from tensorflow.keras.callbacks import ReduceLROnPlateau
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Input
from tensorflow.keras.layers import Activation
from tensorflow.keras.models import load_model
from tensorflow.keras.models import Sequential
from tensorflow.keras.models import Model
from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import roc_auc_score

In [73]: # Point 2
def create_model():
    model = tf.keras.models.Sequential([tf.keras.layers.Dense(2, input_shape=(2,)), kernel_initializer=tf.keras.initializers.RandomUniform(minval=0, maxval=1, seed=42),
    tf.keras.layers.Dense(2, activation='tanh', kernel_initializer=tf.keras.initializers.RandomUniform(minval=0, maxval=1, seed=42),
    tf.keras.layers.Dense(2, activation='tanh', kernel_initializer=tf.keras.initializers.RandomUniform(minval=0, maxval=1, seed=42),
    tf.keras.layers.Dense(2, activation='tanh', kernel_initializer=tf.keras.initializers.RandomUniform(minval=0, maxval=1, seed=42),
    tf.keras.layers.Dense(2, activation='tanh', kernel_initializer=tf.keras.initializers.RandomUniform(minval=0, maxval=1, seed=42),
    tf.keras.layers.Dense(2, activation='sigmoid', name='output_layer')])

In [74]: # Point 3-A Micro F1 Score
class perf_metrics(Callback):
    def on_train_begin(self, logs={}):
        self.f1score = []
        self.precision = []
        self.recall = []

    def on_epoch_end(self, epoch, logs={}):
        y_pred = self.model.predict(X_test).round()
        y_actual = y_test
        val_f1_temp = f1_score(y_actual,y_pred)
        val_precision_temp = precision_score(y_actual,y_pred)
        val_recall_temp = recall_score(y_actual,y_pred)
        self.f1score.append(val_f1_temp)
        self.precision.append(val_precision_temp)
        self.recall.append(val_recall_temp)
        print ("F1 Score : %f, Precision : %f, Recall : %f" % (val_f1_temp, val_precision_temp, val_recall_temp))

    performance_metrics = perf_metrics()
```

```
In [75]: # Point 3-AUC
def scheduler(Callback):
    def __init__(self, training_data, validation_data):
        self.x_train = training_data[0]
        self.y_train = training_data[1]
        self.x_test = validation_data[0]
        self.y_test = validation_data[1]

    def on_epoch_end(self, epoch, logs={}):
        y_pred_train = self.model.predict_proba(self.x_train)
        roc_train = roc_auc_score(self.y_train, y_pred_train)
        y_pred_test = self.model.predict_proba(self.x_test)
        roc_test = roc_auc_score(self.y_test, y_pred_test)
        print ("ROC_AUC_train : %s, ROC_AUC_test : %s" % (str(round(roc_train,4)),str(round(roc_test,4))))
        return roc_auc_score(training_data(X_train, y_train),validation_data(X_test, y_test))

In [76]: filepath = "/content/drive/MyDrive/Calibacks/Model_Save/*Model_1*.weights-(epoch:620)-(val_accuracy:.4f).hdf5"
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')
```

```
In [77]: # Point 5-Condition1
reduce_lr_accuracy = ReduceLROnPlateau(monitor='val_accuracy', Factor=0.9, patience=1, verbose=1, mode='auto')
```

```
In [78]: # Point 6
class TerminateNaN(Callback):
    def on_epoch_end(self, epoch, logs={}):
        loss = logs.get('loss')
        if loss is not None:
            if np.isnan(loss) or np.isinf(loss):
                print("Invalid loss and terminated at epoch {}".format(epoch))
                self.model.stop_training = True
        terminateNaN = TerminateNaN()
```

```
In [79]: # Point 7
earlystop = EarlyStopping(monitor='val_accuracy', patience=2, verbose=1, mode='max')
```

```
In [80]: # Point 8
log_dir = "/content/drive/MyDrive/Calibacks/logs/fit/*Model_1*" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_graph=True, write_grads=True)

WARNING:tensorflow:write_grads will be ignored in TensorFlow 2.0 for the 'TensorBoard' Callback.
```

```
In [81]: # Point 5-Condition2
def scheduler(lr,lr):
    # print (type(lr))
    if (epoch-1)%3==0:
        # print (type(lr)*0.95)
        return lr * 0.95
    else:
        return lr

lr_scheduler = LearningRateScheduler(scheduler, verbose=1)
```

```
In [82]: model_1 = create_model()
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=False, name='SGD')
model_1.compile(optimizer=optimizer,loss='binary_crossentropy',metrics=['accuracy'])
calibacks_list = [performance_metrics,roc,checkpoint,reduce_lr_accuracy,lr_scheduler,terminateNaN,earlystop,tensorboard_callback]
model_1.fit(x=train,y=y_train,epochs=20,validation_data=(X_test, y_test),callbacks=calibacks_list)
```

```
Epoch 1/20
Epoch 00001: LearningRateScheduler reducing learning rate to 0.009999999776482582.
3/419 [.....] - ETA: 27s - loss: 0.7776 - accuracy: 0.4792 WARNING:tensorflow:Callback method 'on_train_batch_end' is slow compared to the batch time (batch time: 0.0024s vs on_train_batch_end time: 0.0223s). Check your callbacks.
419/419 [=====] - 1s 2ms/step - loss: 0.6902 - accuracy: 0.5194 - val_loss: 0.6920 - val_accuracy: 0.5980
F1 Score : 0.641102, Precision : 0.523938, Recall : 0.825758
ROC_AUC train : 0.5867, ROC_AUC test : 0.4937
```

```
Epoch 00001: val_accuracy improved from -inf to 0.49742, saving model to /content/drive/MyDrive/Calibacks/model_save/model_1_weights-01-0.4974.hdf5
Epoch 2/20
Epoch 00002: LearningRateScheduler reducing learning rate to 0.009999999776482582.
419/419 [=====] - 1s 2ms/step - loss: 0.6888 - accuracy: 0.5234 - val_loss: 0.6884 - val_accuracy: 0.5135
F1 Score : 0.581192, Precision : 0.518186, Recall : 0.675152
ROC_AUC train : 0.5685, ROC_AUC test : 0.4937
```

```
Epoch 00003: LearningRateScheduler reducing learning rate to 0.009999999776482582.
419/419 [=====] - 1s 2ms/step - loss: 0.6881 - accuracy: 0.5454 - val_loss: 0.6854 - val_accuracy: 0.5395
F1 Score : 0.646816, Precision : 0.524439, Recall : 0.840363
ROC_AUC train : 0.5887, ROC_AUC test : 0.575
```

```
Epoch 00004: val_accuracy improved from 0.5348 to 0.53773, saving model to /content/drive/MyDrive/Calibacks/model_save/model_1_weights-03-0.5373.hdf5
Epoch 4/20
Epoch 00004: LearningRateScheduler reducing learning rate to 0.0094999997648258196.
419/419 [=====] - 1s 2ms/step - loss: 0.6821 - accuracy: 0.5454 - val_loss: 0.6853 - val_accuracy: 0.5377
F1 Score : 0.641102, Precision : 0.523938, Recall : 0.825758
ROC_AUC train : 0.5867, ROC_AUC test : 0.4937
```

```
Epoch 00004: val_accuracy improved from 0.5348 to 0.53773, saving model to /content/drive/MyDrive/Calibacks/model_save/model_1_weights-04-0.5373.hdf5
Epoch 5/20
Epoch 00005: LearningRateScheduler reducing learning rate to 0.0094999997648258196.
419/419 [=====] - 1s 2ms/step - loss: 0.6852 - accuracy: 0.5392 - val_loss: 0.6854 - val_accuracy: 0.5112
F1 Score : 0.575526, Precision : 0.508685, Recall : 0.662727
ROC_AUC train : 0.5648, ROC_AUC test : 0.5084
```

```
Epoch 00005: val_accuracy did not improve from 0.53773
Epoch 00005: ReduceLROnPlateau reducing learning rate to 0.008549999725873777.
Epoch 6/20
Epoch 00006: LearningRateScheduler reducing learning rate to 0.008122499108863934.
419/419 [=====] - 1s 2ms/step - loss: 0.6821 - accuracy: 0.5454 - val_loss: 0.6834 - val_accuracy: 0.5395
F1 Score : 0.646816, Precision : 0.524439, Recall : 0.840363
ROC_AUC train : 0.5887, ROC_AUC test : 0.575
```

```
Epoch 00006: val_accuracy improved from 0.53773 to 0.53955, saving model to /content/drive/MyDrive/Calibacks/model_save/model_1_weights-06-0.5395.hdf5
Epoch 7/20
Epoch 00007: LearningRateScheduler reducing learning rate to 0.008122499108863934.
419/419 [=====] - 1s 2ms/step - loss: 0.6819 - accuracy: 0.5428 - val_loss: 0.6836 - val_accuracy: 0.5391
F1 Score : 0.646816, Precision : 0.524439, Recall : 0.840363
ROC_AUC train : 0.5946, ROC_AUC test : 0.5858
```

```
Epoch 00007: val_accuracy did not improve from 0.53955
Epoch 00007: ReduceLROnPlateau reducing learning rate to 0.00731624919077754.
Epoch 8/20
Epoch 00008: LearningRateScheduler reducing learning rate to 0.00731624919077754.
419/419 [=====] - 1s 2ms/step - loss: 0.6812 - accuracy: 0.5481 - val_loss: 0.6822 - val_accuracy: 0.5409
F1 Score : 0.650929, Precision : 0.524439, Recall : 0.839118
ROC_AUC train : 0.5996, ROC_AUC test : 0.5984
```

```
Epoch 00008: val_accuracy improved from 0.53955 to 0.54091, saving model to /content/drive/MyDrive/Calibacks/model_save/model_1_weights-08-0.5409.hdf5
Epoch 9/20
Epoch 00009: LearningRateScheduler reducing learning rate to 0.0069449997648258196.
419/419 [=====] - 1s 2ms/step - loss: 0.6804 - accuracy: 0.5484 - val_loss: 0.6816 - val_accuracy: 0.5414
F1 Score : 0.660884, Precision : 0.524439, Recall : 0.839118
ROC_AUC train : 0.6083, ROC_AUC test : 0.598
```

```
Epoch 00009: val_accuracy improved from 0.54091 to 0.54136, saving model to /content/drive/MyDrive/Calibacks/model_save/model_1_weights-09-0.5414.hdf5
Epoch 10/20
Epoch 00010: LearningRateScheduler reducing learning rate to 0.0069449997648258196.
419/419 [=====] - 1s 2ms/step - loss: 0.6801 - accuracy: 0.5525 - val_loss: 0.6806 - val_accuracy: 0.5418
F1 Score : 0.660825, Precision : 0.524439, Recall : 0.839118
ROC_AUC train : 0.6085, ROC_AUC test : 0.598
```

```
Epoch 00010: val_accuracy improved from 0.54136 to 0.54182, saving model to /content/drive/MyDrive/Calibacks/model_save/model_1_weights-10-0.5418.hdf5
Epoch 11/20
Epoch 00011: LearningRateScheduler reducing learning rate to 0.0069449997648258196.
419/419 [=====] - 1s 2ms/step - loss: 0.6804 - accuracy: 0.5516 - val_loss: 0.6829 - val_accuracy: 0.5397
F1 Score : 0.655952, Precision : 0.523949, Recall : 0.868485
ROC_AUC train : 0.6, ROC_AUC test : 0.5979
```

```
Epoch 00011: val_accuracy did not improve from 0.54182
Epoch 00011: ReduceLROnPlateau reducing learning rate to 0.006250623832357515.
Epoch 12/20
Epoch 00012: LearningRateScheduler reducing learning rate to 0.00593758023715807.
419/419 [=====] - 1s 2ms/step - loss: 0.6789 - accuracy: 0.5576 - val_loss: 0.6800 - val_accuracy: 0.5398
F1 Score : 0.656728, Precision : 0.523707, Recall : 0.880283
ROC_AUC train : 0.6018, ROC_AUC test : 0.5999
```

```
Epoch 00012: val_accuracy did not improve from 0.54182
Epoch 00012: ReduceLROnPlateau reducing learning rate to 0.005344974987844931.
Epoch 00012: early stopping
tensorflow.python.keras.callbacks.History at 0x73862358960
```

```
In [83]: model_1.summary()

Model: "sequential_18"
Layer (type) Output Shape Param #
-----
layer1 (Dense) (None, 2) 6
layer2 (Dense) (None, 2) 6
layer3 (Dense) (None, 2) 6
layer4 (Dense) (None, 2) 6
layer5 (Dense) (None, 2) 6
output_layer (Dense) (None, 1) 3
Total params: 33
Trainable params: 33
Non-trainable params: 0
```

```
In [84]: !reload_ext tensorboard
!rm -rf ./logs/
!tensorboard --logdir=/content/drive/MyDrive/Calibacks/logs/fit/

Output hidden; open in https://colab.research.google.com to view.
```

```
Model-2
1. Use relu as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use he_uniform() as initializer.
3. Analyze your output and training process.
```

```
In [86]: def create_model2():
    return tf.keras.models.Sequential([tf.keras.layers.Dense(2, input_shape=(2,)), kernel_initializer=tf.keras.initializers.RandomUniform(minval=0, maxval=1, seed=42),
    tf.keras.layers.Dense(2, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform(), name='layer1'),
    tf.keras.layers.Dense(2, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform(), name='layer2'),
    tf.keras.layers.Dense(2, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform(), name='layer3'),
    tf.keras.layers.Dense(2, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform(), name='layer4'),
    tf.keras.layers.Dense(2, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform(), name='layer5'),
    tf.keras.layers.Dense(1, activation='sigmoid', name='output_layer')])

In [88]: filepath = "/content/drive/MyDrive/Calibacks/Model_Save/*Model_2*.weights-(epoch:620)-(val_accuracy:.4f).hdf5"
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')
log_dir = "/content/drive/MyDrive/Calibacks/logs/fit/*Model_2*" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_graph=True, write_grads=True)

WARNING:tensorflow:write_grads will be ignored in TensorFlow 2.0 for the 'TensorBoard' Callback.
```

```
In [89]: model_2 = create_model2()
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=False, name='SGD')
model_2.compile(optimizer=optimizer,loss='binary_crossentropy',metrics=['accuracy'])
calibacks_list = [checkpoint,tensorboard_callback,earlystop,reduce_lr_accuracy,lr_scheduler,terminateNaN,performance_metrics,roc]
model_2.fit(x=train,y=y_train,epochs=20,validation_data=(X_test, y_test),callbacks=calibacks_list)
```

```
Epoch 1/20
Epoch 00001: LearningRateScheduler reducing learning rate to 0.009999999776482582.
3/419 [.....] - ETA: 28s - loss: 0.8854 - accuracy: 0.5478 WARNING:tensorflow:Callback method 'on_train_batch_end' is slow compared to the batch time (batch time: 0.0024s vs on_train_batch_end time: 0.0213s). Check your callbacks.
419/419 [=====] - 1s 2ms/step - loss: 0.6804 - accuracy: 0.5451 - val_loss: 0.6854 - val_accuracy: 0.5324
F1 Score : 0.630834, Precision : 0.52145, Recall : 0.796677
ROC_AUC train : 0.5515, ROC_AUC test : 0.5388
```

```
Epoch 2/20
Epoch 00002: LearningRateScheduler reducing learning rate to 0.009999999776482582.
419/419 [=====] - 1s 2ms/step - loss: 0.6789 - accuracy: 0.5552 - val_loss: 0.6722 - val_accuracy: 0.5791
F1 Score : 0.629806, Precision : 0.524426, Recall : 0.718636
ROC_AUC train : 0.6020, ROC_AUC test : 0.6062
```

```
Epoch 00003: LearningRateScheduler reducing learning rate to 0.0094999997648258196.
419/419 [=====] - 1s 2ms/step - loss: 0.6799 - accuracy: 0.5577 - val_loss: 0.6748 - val_accuracy: 0.5698
F1 Score : 0.652807, Precision : 0.547263, Recall : 0.888788
ROC_AUC train : 0.5915, ROC_AUC test : 0.5642
```

```
Epoch 00004: LearningRateScheduler reducing learning rate to 0.008549999725873777.
419/419 [=====] - 1s 2ms/step - loss: 0.6804 - accuracy: 0.5493 - val_loss: 0.6854 - val_accuracy: 0.5380
F1 Score : 0.642816, Precision : 0.524856, Recall : 0.826485
ROC_AUC train : 0.5807, ROC_AUC test : 0.5392
```

```
Epoch 00005: val_accuracy did not improve from 0.54273
Epoch 00005: ReduceLROnPlateau reducing learning rate to 0.007676499959109213.
F1 Score : 0.642816, Precision : 0.524856, Recall : 0.826485
ROC_AUC train : 0.5807, ROC_AUC test : 0.5392
```

```
Epoch 00005: early stopping
tensorflow.python.keras.callbacks.History at 0x73862358960
```

```
In [90]: model_2.summary()

Model: "sequential_19"
Layer (type) Output Shape Param #
-----
layer1 (Dense) (None, 2) 6
layer2 (Dense) (None, 2) 6
layer3 (Dense) (None, 2) 6
layer4 (Dense) (None, 2) 6
layer5 (Dense) (None, 2) 6
output_layer (Dense) (None, 1) 3
Total params: 33
Trainable params: 33
Non-trainable params: 0
```

```
In [91]: !reload_ext tensorboard
!rm -rf ./logs/
!tensorboard --logdir=/content/drive/MyDrive/Calibacks/logs/fit/

Output hidden; open in https://colab.research.google.com to view.
```

```
Model-3
1. Use relu as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use he_uniform() as initializer.
3. Analyze your output and training process.
```

```
In [92]: def create_model3():
    return tf.keras.models.Sequential([tf.keras.layers.Dense(2, input_shape=(2,)), kernel_initializer=tf.keras.initializers.HeUniform(), name='layer1'),
    tf.keras.layers.Dense(2, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform(), name='layer2'),
    tf.keras.layers.Dense(2, activation='tanh', kernel_initializer=tf.keras.initializers.HeUniform(), name='layer3'),
    tf.keras.layers.Dense(2, activation='tanh', kernel_initializer=tf.keras.initializers.HeUniform(), name='layer4'),
    tf.keras.layers.Dense(2, activation='sigmoid', name='output_layer')])

In [93]: filepath = "/content/drive/MyDrive/Calibacks/Model_Save/*Model_3*.weights-(epoch:620)-(val_accuracy:.4f).hdf5"
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')
log_dir = "/content/drive/MyDrive/Calibacks/logs/fit/*Model_3*" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_graph=True, write_grads=True)

WARNING:tensorflow:write_grads will be ignored in TensorFlow 2.0 for the 'TensorBoard' Callback.
```

```
In [94]: model_3 = create_model3()
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=False, name='SGD')
model_3.compile(optimizer=optimizer,loss='binary_crossentropy',metrics=['accuracy'])
calibacks_list = [checkpoint,tensorboard_callback,earlystop,reduce_lr_accuracy,lr_scheduler,terminateNaN,performance_metrics,roc]
model_3.fit(x=train,y=y_train,epochs=20,validation_data=(X_test, y_test),callbacks=calibacks_list)
```

```
Epoch 1/20
Epoch 00001: LearningRateScheduler reducing learning rate to 0.009999999776482582.
3/419 [.....] - ETA: 28s - loss: 0.8854 - accuracy: 0.5478 WARNING:tensorflow:Callback method 'on_train_batch_end' is slow compared to the batch time (batch time: 0.0024s vs on_train_batch_end time: 0.0213s). Check your callbacks.
419/419 [=====] - 1s 2ms/step - loss: 0.6804 - accuracy: 0.5451 - val_loss: 0.6854 - val_accuracy: 0.5324
F1 Score : 0.630834, Precision : 0.52145, Recall : 0.796677
ROC_AUC train : 0.5515, ROC_AUC test : 0.5388
```

```
Epoch 2/20
Epoch 00002: LearningRateScheduler reducing learning rate to 0.009999999776482582.
419/419 [=====] - 1s 2ms/step - loss: 0.6789 - accuracy: 0.5552 - val_loss: 0.6722 - val_accuracy: 0.5791
F1 Score : 0.629806, Precision : 0.524426, Recall : 0.718636
ROC_AUC train : 0.6020, ROC_AUC test : 0.6062
```

```
Epoch 00003: LearningRateScheduler reducing learning rate to 0.0094999997648258196.
419/419 [=====] - 1s 2ms/step - loss: 0.6799 - accuracy: 0.5577 - val_loss: 0.6748 - val_accuracy: 0.5698
F1 Score : 0.652807, Precision : 0.547263, Recall : 0.888788
ROC_AUC train : 0.5915, ROC_AUC test : 0.5642
```

```
Epoch 00004: LearningRateScheduler reducing learning rate to 0.008549999725873777.
419/419 [=====] - 1s 2ms/step - loss: 0.6804 - accuracy: 0.5493 - val_loss: 0.6854 - val_accuracy: 0.5380
F1 Score : 0.642816, Precision : 0.524856, Recall : 0.826485
ROC_AUC train : 0.5807, ROC_AUC test : 0.5392
```

```
Epoch 00005: val_accuracy did not improve from 0.54273
Epoch 00005: ReduceLROnPlateau reducing learning rate to 0.007676499959109213.
F1 Score : 0.642816, Precision : 0.524856, Recall : 0.826485
ROC_AUC train : 0.5807, ROC_AUC test : 0.5392
```

```
Epoch 00005: early stopping
tensorflow.python.keras.callbacks.History at 0x73862358960
```

```
In [95]: model_3.summary()

Model: "sequential_21"
Layer (type) Output Shape Param #
-----
layer1 (Dense) (None, 2) 6
layer2 (Dense) (None, 2) 6
layer3 (Dense) (None, 2) 6
layer4 (Dense) (None, 2) 6
layer5 (Dense) (None, 2) 6
output_layer (Dense) (None, 1) 3
Total params: 33
Trainable params: 33
Non-trainable params: 0
```

```
In [96]: !reload_ext tensorboard
!rm -rf ./logs/
!tensorboard --logdir=/content/drive/MyDrive/Calibacks/logs/fit/

Output hidden; open in https://colab.research.google.com to view.
```

```
Model-4
1. Try with any values to get better accuracy/F1 score.
```

```
In [96]: def create_model4():
    return tf.keras.models.Sequential([tf.keras.layers.Dense(2, input_shape=(2,)), kernel_initializer=tf.keras.initializers.GlorotNormal(seed=None), name='layer1'),
    tf.keras.layers.Dense(2, activation='relu', kernel_initializer=tf.keras.initializers.GlorotNormal(seed=None), name='layer2'),
    tf.keras.layers.Dense(2, activation='tanh', kernel_initializer=tf.keras.initializers.GlorotNormal(seed=None), name='layer3'),
    tf.keras.layers.Dense(2, activation='tanh', kernel_initializer=tf.keras.initializers.GlorotNormal(seed=None), name='layer4'),
    tf.keras.layers.Dense(2, activation='sigmoid', name='output_layer')])

In [97]: filepath = "/content/drive/MyDrive/Calibacks/Model_Save/*Model_4*.weights-(epoch:620)-(val_accuracy:.4f).hdf5"
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')
log_dir = "/content/drive/MyDrive/Calibacks/logs/fit/*Model_4*" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_graph=True, write_grads=True)

WARNING:tensorflow:write_grads will be ignored in TensorFlow 2.0 for the 'TensorBoard' Callback.
```

```
In [98]: model_4 = create_model4()
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=False, name='SGD')
model_4.compile(optimizer=optimizer,loss='binary_crossentropy',metrics=['accuracy'])
calibacks_list = [checkpoint,tensorboard_callback,earlystop,reduce_lr_accuracy,lr_scheduler,terminateNaN,performance_metrics,roc]
model_4.fit(x=train,y=y_train,epochs=20,validation_data=(X_test, y_test),callbacks=calibacks_list)
```

```
Epoch 1/20
Epoch 00001: LearningRateScheduler reducing learning rate to 0.009999999776482582.
3/419 [.....] - ETA: 28s - loss: 0.8854 - accuracy: 0.5478 WARNING:tensorflow:Callback method 'on_train_batch_end' is slow compared to the batch time (batch time: 0.0024s vs on_train_batch_end time: 0.0213s). Check your callbacks.
419/419 [=====] - 1s 2ms/step - loss: 0.6804 - accuracy: 0.5451 - val_loss: 0.6854 - val_accuracy: 0.5324
F1 Score : 0.666602, Precision : 0.500000, Recall : 1.000000
ROC_AUC train : 0.5306, ROC_AUC test : 0.5213
```

```
Epoch 2/20
Epoch 00002: LearningRateScheduler reducing learning rate to 0.009999999776482582.
419/419 [=====] - 1s 2ms/step - loss: 0.6789 - accuracy: 0.5523 - val_loss: 0.6843 - val_accuracy: 0.5744
F1 Score : 0.659806, Precision : 0.507599, Recall : 0.819394
ROC_AUC train : 0.5958, ROC_AUC test : 0.5827
```

```
Epoch 00003: LearningRateScheduler reducing learning rate to 0.0094999997648258196.
419/419 [=====] - 1s 2ms/step - loss: 0.6742 - accuracy: 0.5954 - val_loss: 0.6652 - val_accuracy: 0.6028
F1 Score : 0.652807, Precision : 0.547263, Recall : 0.770989
ROC_AUC train : 0.6115, ROC_AUC test : 0.6028
```

```
Epoch 00004: LearningRateScheduler reducing learning rate to 0.008549999725873777.
419/419 [=====] - 1s 2ms/step - loss: 0.6804 - accuracy: 0.5493 - val_loss: 0.6854 - val_accuracy: 0.5380
F1 Score : 0.642816, Precision : 0.524856, Recall : 0.826485
ROC_AUC train : 0.5807, ROC_AUC test : 0.5392
```

```
Epoch 00005: val_accuracy did not improve from 0.60258
Epoch 00005: ReduceLROnPlateau reducing learning rate to 0.007676499959109213.
F1 Score : 0.642816, Precision : 0.524856, Recall : 0.826485
ROC_AUC train : 0.5807, ROC_AUC test : 0.5392
```

```
Epoch 00005: early stopping
tensorflow.python.keras.callbacks.History at 0x73862358960
```

```
In [99]: model_4.summary()

Model: "sequential_26"
Layer (type) Output Shape Param #
-----
layer1 (Dense) (None, 2) 6
layer2 (Dense) (None, 2) 6
layer3 (Dense) (None, 2) 6
layer4 (Dense) (None, 2) 6
layer5 (Dense) (None, 2) 6
output_layer (Dense) (None, 1) 3
Total params: 33
Trainable params: 33
Non-trainable params: 0
```

```
In [100]: !reload_ext tensorboard
!rm -rf ./logs/
!tensorboard --logdir=/content/drive/MyDrive/Calibacks/logs/fit/

Output hidden; open in https://colab.research.google.com to view.
```

```
In [ ]: 
```